Wine Quality Prediction AWS Spark Application:

Pa2Winepred: The creation of a Python application using the PySpark interface is the goal of this project.

An Elastic MapReduce (EMR) cluster from Amazon Web Services (AWS) is used to host the application. The main goal is to use publicly available data to train a machine learning model in parallel on EC2 instances to predict wine quality. The trained algorithm is then used to forecast the wine's quality. The trained machine learning model's container image is made using Docker, which expedites the deployment procedure.

Link for GitHub:

https://github.com/Sanjeev-Prem/Cloud_Assi2

Link for Docker:

https://hub.docker.com/repository/docker/sg2656/predequlity

Steps for the Execution for Wine Quality Prediction AWS Spark Application:

1.Create a Key-pair for the EMR Cluster :go to EC2/Network/Key-pairs

Use the format of .pem and download the keypair

Created key pair as: Prem.pem

2.Create an S3 bucket

Created an S3 bucket in aws: Buicketwine

3.Then go to EMR console and create EMR cluster

4. Creating the spark in the AWS instance by using EMR console:

Creating the spark cluster by using the EMR console, and create the 4 instances:

Name and application:

Name:pa2winepqmahi

Amazon EMR release: emr-5.33.0

Application bundle: Hadoop 2.10.1, Spark 2.4.7, Zippeline 0.9.0, and Yarn

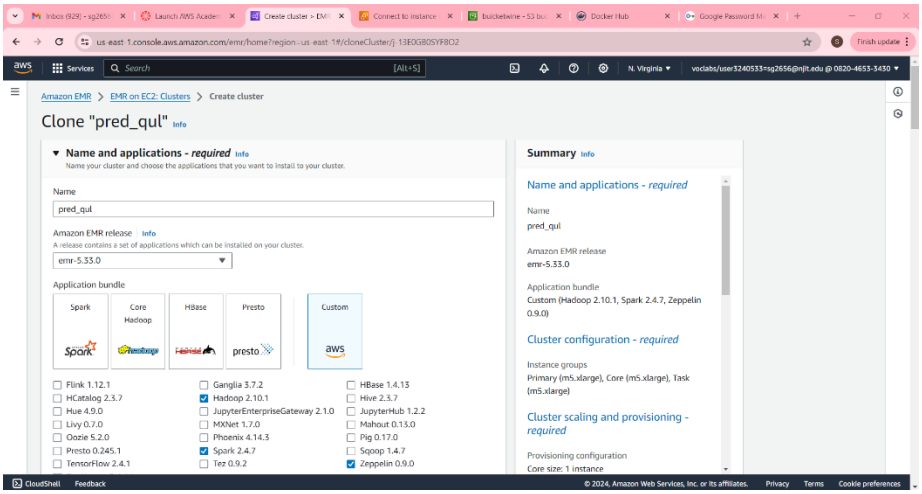

Note: Here it says Clone “pred_qul” as I have cloned the previous configuration instead of creating from scratch to save time.

Cluster Configuration:

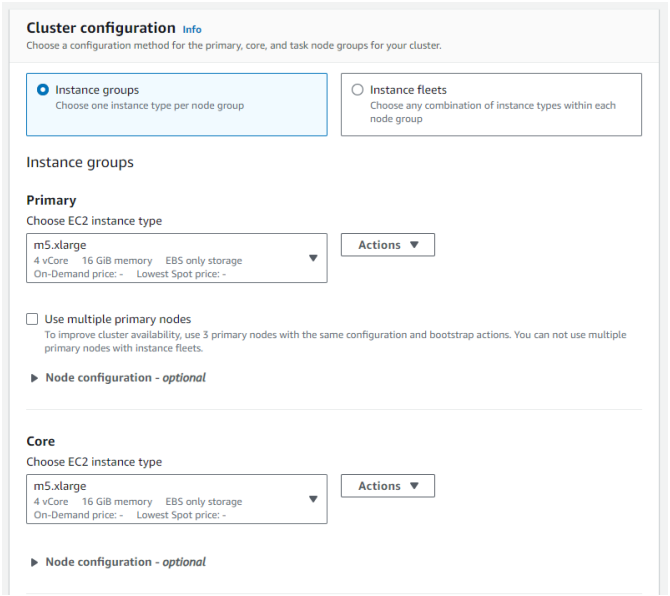

Cluster Scaling and provisioning:

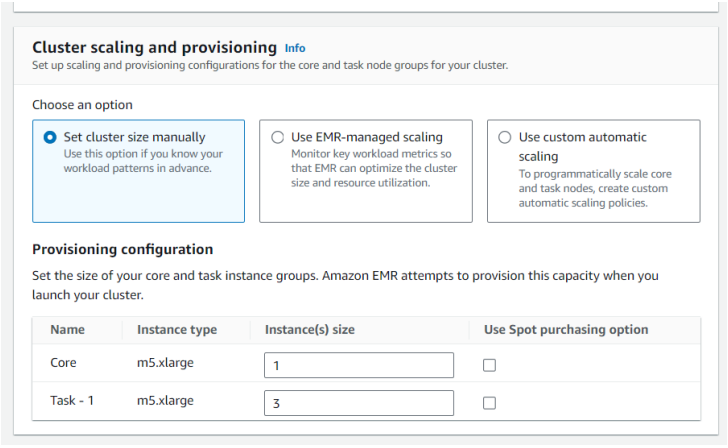

Networking & Cluster Termination:

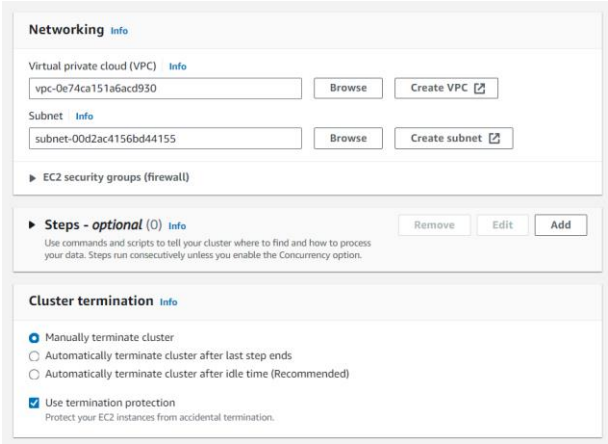

Security Configuration and EC2 Key pair & Identity and access management(IAM) roles:

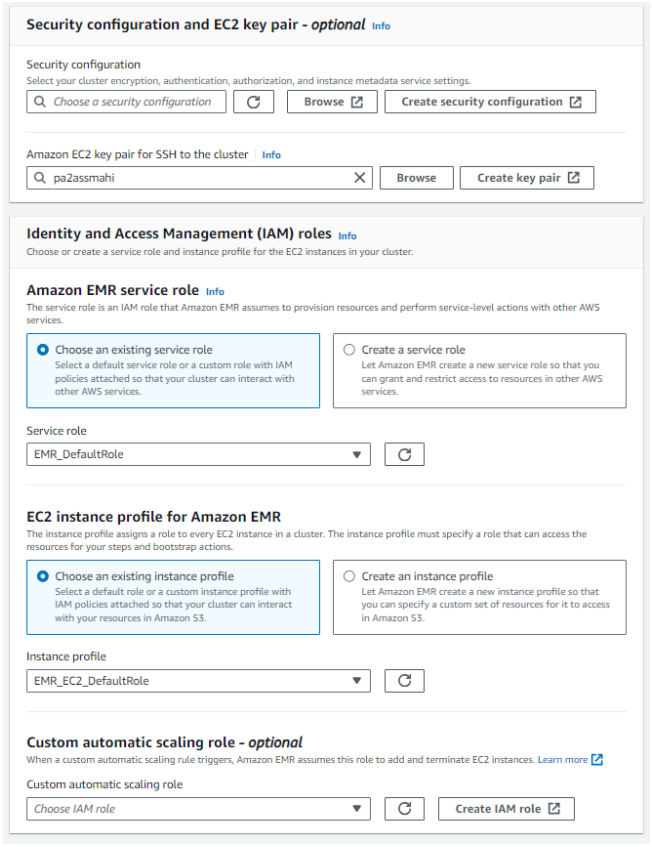

We can follow above steps for creating the EMR cluster for the instances

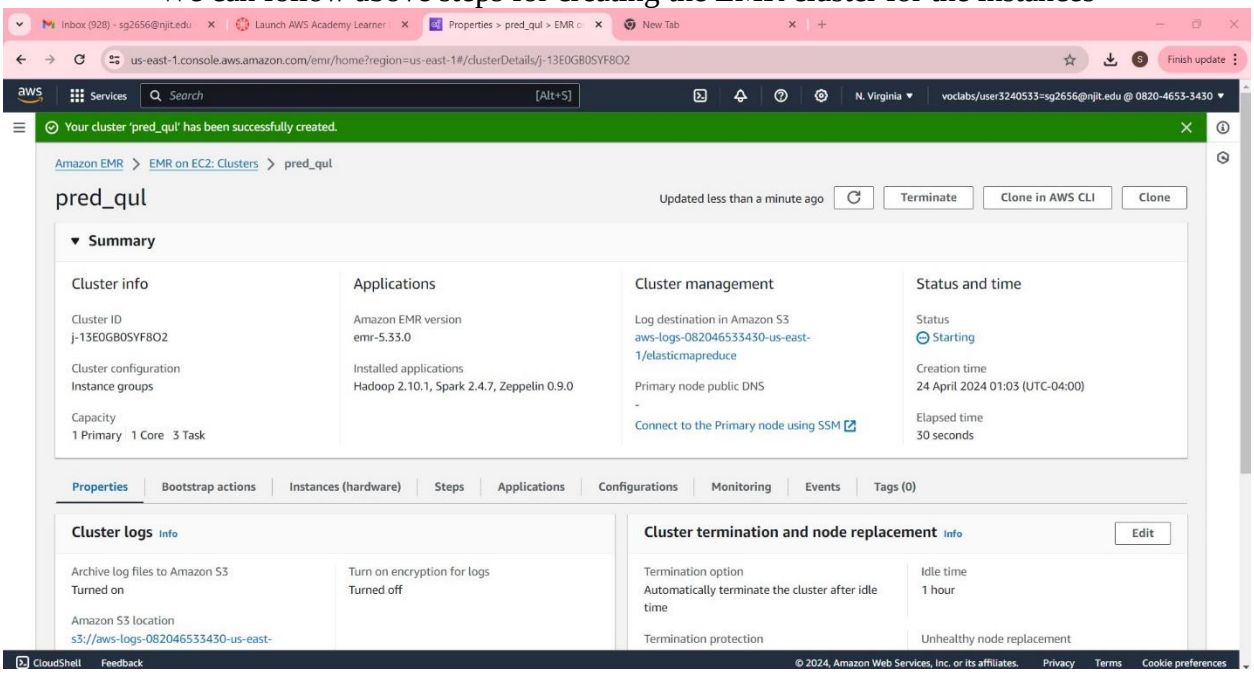

5. We are now parallelly training an ML model on an EC2 instance in a spark cluster without Docker:

1.Now the cluster will accept the tasks to run the ML model

Need to connect the Master instance in the Terminal:

ssh -i "Prem.pem" ec2-user@ec2-54-205-26-60.compute-1.amazonaws.com

and it is successfully login.

2.After the login of Master instance then change the root by using

Sudo su

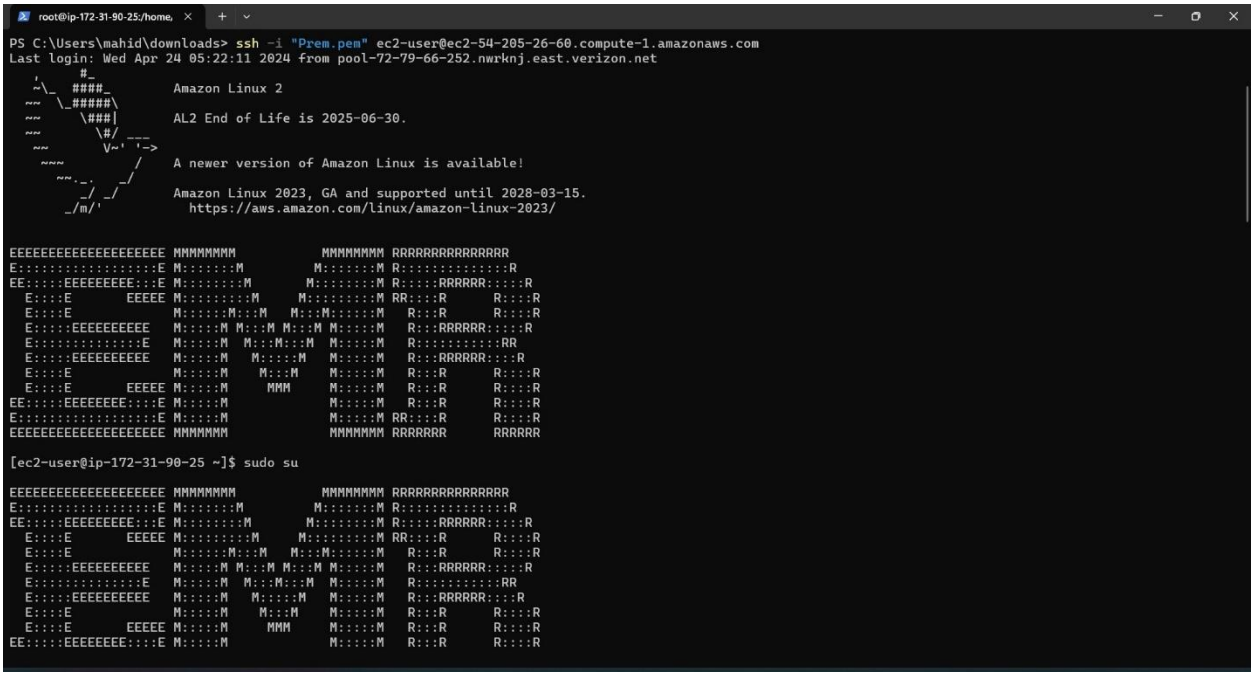

3. Submit the task by the command:

spark-submit s3://buicketwine/ quilityprediction.py

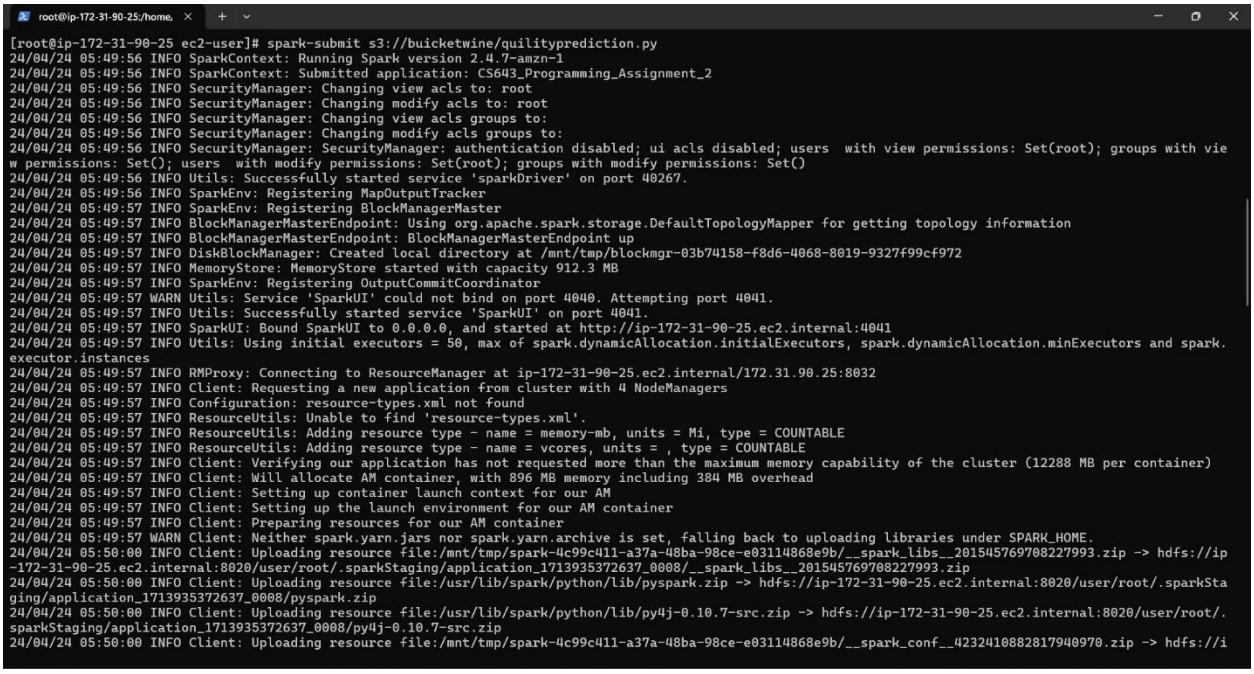
root@ip-172-31-90-25:/home/
[root@ip-172-31-90-25 ec2-user]# spark-submit s3://buicketwine/quilityprediction.py
24/04/24 05:49:56 INFO SparkContext: Running Spark version 2.4.7-amzn-1
24/04/24 05:49:56 INFO SparkContext: Submitted application: CS643_Programming_Assignment_2
24/04/24 05:49:56 INFO SecurityManager: Changing view acls to: root
24/04/24 05:49:56 INFO SecurityManager: Changing modify acls to: root
24/04/24 05:49:56 INFO SecurityManager: Changing view acls groups to:
24/04/24 05:49:56 INFO SecurityManager: Changing modify acls groups to:
24/04/24 05:49:56 INFO SecurityManager: SecurityManager: authentication disabled; ui acls disabled; users with view permissions: Set(root); groups with vie
w permissions: Set(); users with modify permissions: Set(root); groups with modify permissions: Set()
24/04/24 05:49:56 INFO Utils: Successfully started service 'sparkDriver' on port 40267.
24/04/24 05:49:56 INFO SparkEnv: Registering MapOutputTracker
24/04/24 05:49:57 INFO SparkEnv: Registering BlockManagerMaster
24/04/24 05:49:57 INFO BlockManagerMasterEndpoint: Using org.apache.spark.storage.DefaultTopologyMapper for getting topology information
24/04/24 05:49:57 INFO BlockManagerMasterEndpoint: BlockManagerMasterEndpoint up
24/04/24 05:49:57 INFO DiskBlockManager: Created local directory at /mnt/tmp/blockmgr-03b74158-f8d6-4068-8019-9327f99cf972
24/04/24 05:49:57 INFO MemoryStore: MemoryStore started with capacity 912.3 MB
24/04/24 05:49:57 INFO SparkEnv: Registering OutputCommitCoordinator
24/04/24 05:49:57 WARN Utils: Service 'SparkUI' could not bind on port 4040. Attempting port 4041.
24/04/24 05:49:57 INFO Utils: Successfully started service 'SparkUI' on port 4041.
24/04/24 05:49:57 INFO SparkUI: Bound SparkUI to 0.0.0.0, and started at http://ip-172-31-90-25.ec2.internal:4041
24/04/24 05:49:57 INFO Utils: Using initial executors = 50, max of spark.dynamicAllocation.initialExecutors, spark.dynamicAllocation.minExecutors and spark.
executor.instances
24/04/24 05:49:57 INFO RMProxy: Connecting to ResourceManager at ip-172-31-90-25.ec2.internal/172.31.90.25:8032
24/04/24 05:49:57 INFO Client: Requesting a new application from cluster with 4 NodeManagers
24/04/24 05:49:57 INFO Configuration: resource-types.xml not found
24/04/24 05:49:57 INFO ResourceUtils: Unable to find 'resource-types.xml'.
24/04/24 05:49:57 INFO ResourceUtils: Adding resource type - name = memory-mb, units = Mi, type = COUNTABLE
24/04/24 05:49:57 INFO ResourceUtils: Adding resource type - name = vcores, units = , type = COUNTABLE
24/04/24 05:49:57 INFO Client: Verifying our application has not requested more than the maximum memory capability of the cluster (12288 MB per container)
24/04/24 05:49:57 INFO Client: Will allocate AM container, with 896 MB memory including 384 MB overhead
24/04/24 05:49:57 INFO Client: Setting up container launch context for our AM
24/04/24 05:49:57 INFO Client: Setting up the launch environment for our AM container
24/04/24 05:49:57 INFO Client: Preparing resources for our AM container
24/04/24 05:49:57 WARN Client: Neither spark.yarn.jars nor spark.yarn.archive is set, falling back to uploading libraries under SPARK_HOME.
24/04/24 05:50:00 INFO Client: Uploading resource file:/mnt/tmp/spark-4c99c411-a37a-48ba-98ce-e03114868e9b/__spark_libs__2015457697082279 93.zip -> hdfs://ip
-172-31-90-25.ec2.internal:8020/user/root/.sparkStaging/application_1713935372637_0008/__spark_libs__2015457697082279 93.zip
24/04/24 05:50:00 INFO Client: Uploading resource file:/usr/lib/spark/python/lib/pyspark.zip -> hdfs://ip-172-31-90-25.ec2.internal:8020/user/root/.sparkSta
ging/application_1713935372637_0008/pyspark.zip
24/04/24 05:50:00 INFO Client: Uploading resource file:/usr/lib/spark/python/lib/py4j-0.10.7-src.zip -> hdfs://ip-172-31-90-25.ec2.internal:8020/user/root/.
sparkStaging/application_1713935372637_0008/py4j-0.10.7-src.zip
24/04/24 05:50:00 INFO Client: Uploading resource file:/mnt/tmp/spark-4c99c411-a37a-48ba-98ce-e03114868e9b/__spark_conf__4232410882817940970.zip -> hdfs://i

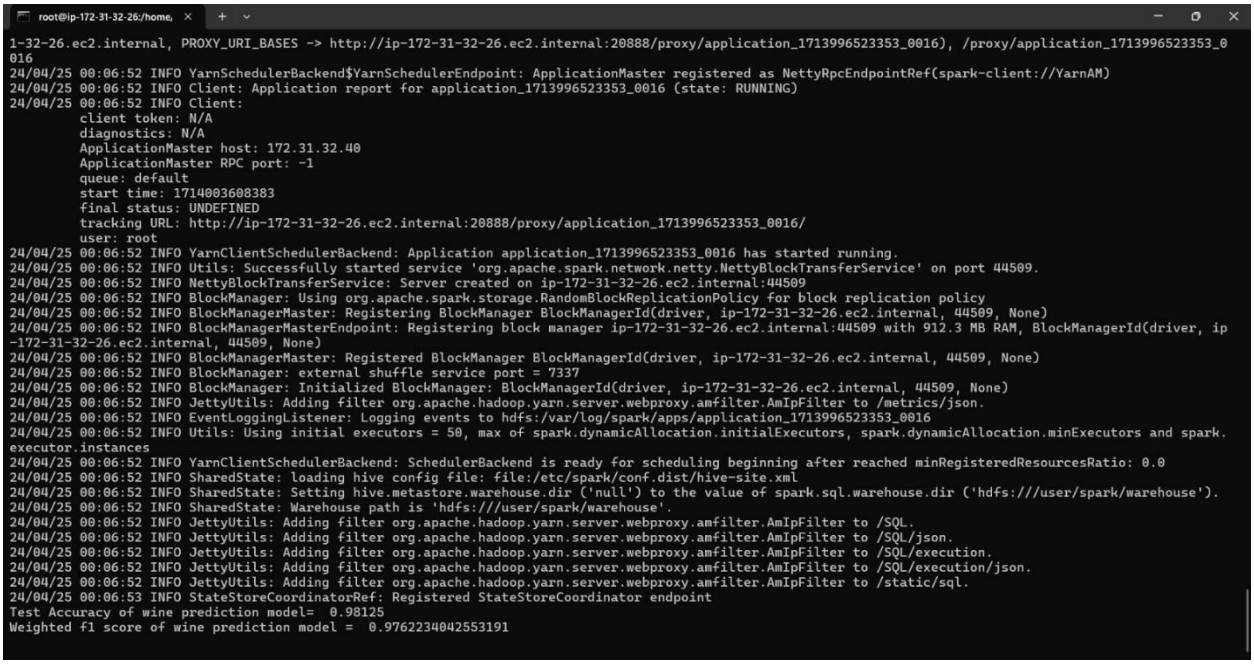
root@ip-172-31-32-26:/home/
1-32-26.ec2.internal, PROXY_URI_BASES -> http://ip-172-31-32-26.ec2.internal:20888/proxy/application_1713996523353_0016), /proxy/application_1713996523353_0
016
24/04/25 00:06:52 INFO YarnSchedulerBackend$YarnSchedulerEndpoint: ApplicationMaster registered as NettyRpcEndpointRef(spark-client://YarnAM)
24/04/25 00:06:52 INFO Client: Application report for application_1713996523353_0016 (state: RUNNING)
24/04/25 00:06:52 INFO Client:
client token: N/A
diagnostics: N/A
ApplicationMaster host: 172.31.32.40
ApplicationMaster RPC port: -1
queue: default
start time: 1714003608383
final status: UNDEFINED
tracking URL: http://ip-172-31-32-26.ec2.internal:20888/proxy/application_1713996523353_0016/
user: root
24/04/25 00:06:52 INFO YarnClientSchedulerBackend: Application application_1713996523353_0016 has started running.
24/04/25 00:06:52 INFO Utils: Successfully started service 'org.apache.spark.network.netty.NettyBlockTransferService' on port 44509.
24/04/25 00:06:52 INFO NettyBlockTransferService: Server created on ip-172-31-32-26.ec2.internal:44509
24/04/25 00:06:52 INFO BlockManager: Using org.apache.spark.storage.RandomBlockReplicationPolicy for block replication policy
24/04/25 00:06:52 INFO BlockManagerMaster: Registering BlockManager BlockManagerId(driver, ip-172-31-32-26.ec2.internal, 44509, None)
24/04/25 00:06:52 INFO BlockManagerMasterEndpoint: Registering block manager ip-172-31-32-26.ec2.internal:44509 with 912.3 MB RAM, BlockManagerId(driver, ip
-172-31-32-26.ec2.internal, 44509, None)
24/04/25 00:06:52 INFO BlockManagerMaster: Registered BlockManager BlockManagerId(driver, ip-172-31-32-26.ec2.internal, 44509, None)
24/04/25 00:06:52 INFO BlockManager: external shuffle service port = 7337
24/04/25 00:06:52 INFO BlockManager: Initialized BlockManager: BlockManagerId(driver, ip-172-31-32-26.ec2.internal, 44509, None)
24/04/25 00:06:52 INFO JettyUtils: Adding filter org.apache.hadoop.yarn.server.webproxy.amfilter.AmIpFilter to /metrics/json.
24/04/25 00:06:52 INFO EventLoggingListener: Logging events to hdfs:/var/log/spark/apps/application_1713996523353_0016
24/04/25 00:06:52 INFO Utils: Using initial executors = 50, max of spark.dynamicAllocation.initialExecutors, spark.dynamicAllocation.minExecutors and spark.
executor.instances
24/04/25 00:06:52 INFO YarnClientSchedulerBackend: SchedulerBackend is ready for scheduling beginning after reached minRegisteredResourcesRatio: 0.0
24/04/25 00:06:52 INFO SharedState: loading hive config file: file:/etc/spark/conf.dist/hive-site.xml
24/04/25 00:06:52 INFO SharedState: Setting hive.metastore.warehouse.dir ('null') to the value of spark.sql.warehouse.dir ('hdfs:///user/spark/warehouse').
24/04/25 00:06:52 INFO SharedState: Warehouse path is 'hdfs:///user/spark/warehouse'.
24/04/25 00:06:52 INFO JettyUtils: Adding filter org.apache.hadoop.yarn.server.webproxy.amfilter.AmIpFilter to /SQL.
24/04/25 00:06:52 INFO JettyUtils: Adding filter org.apache.hadoop.yarn.server.webproxy.amfilter.AmIpFilter to /SQL/json.
24/04/25 00:06:52 INFO JettyUtils: Adding filter org.apache.hadoop.yarn.server.webproxy.amfilter.AmIpFilter to /SQL/execution.
24/04/25 00:06:52 INFO JettyUtils: Adding filter org.apache.hadoop.yarn.server.webproxy.amfilter.AmIpFilter to /SQL/execution/json.
24/04/25 00:06:52 INFO JettyUtils: Adding filter org.apache.hadoop.yarn.server.webproxy.amfilter.AmIpFilter to /static/sql.
24/04/25 00:06:53 INFO StateStoreCoordinatorRef: Registered StateStoreCoordinator endpoint
Test Accuracy of wine prediction model= 0.98125
Weighted f1 score of wine prediction model = 0.9762234042553191

4.Then you can find the trace status for the above tasks, The status is succeed then there is a creation of test.model in the s3 bucket s3://buicketwine

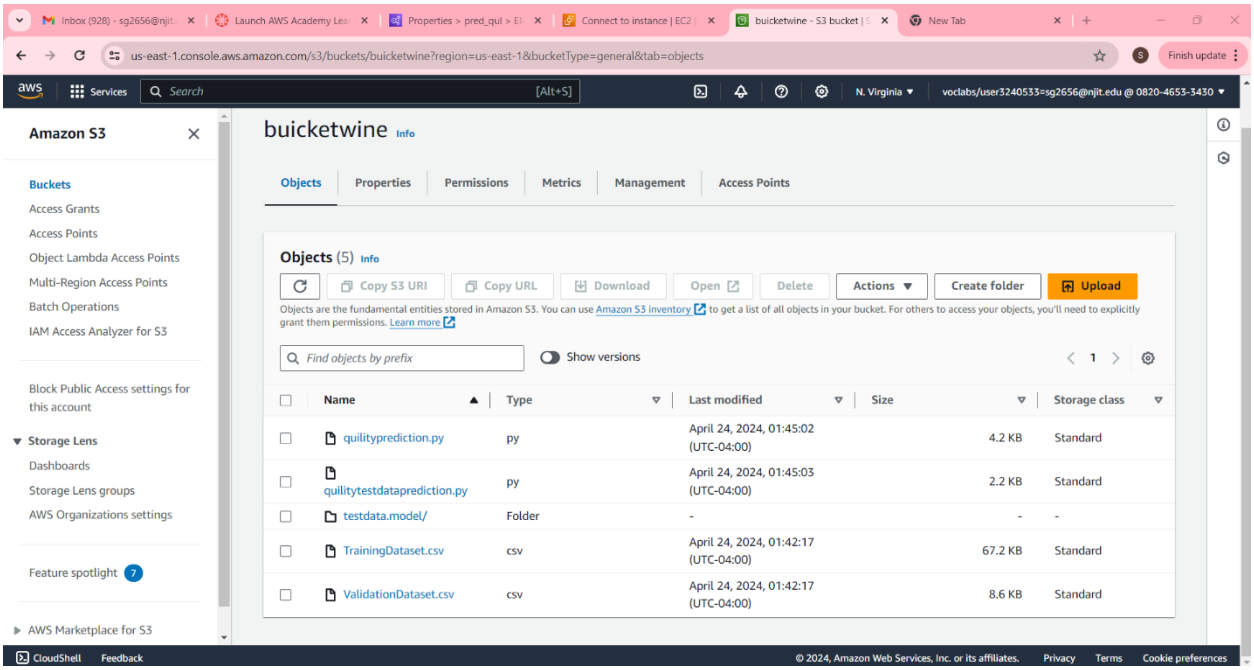

6.Now we are running ML model using the Docker:

1.Create an docker account and sign up.

2. After the successful login then download and setup the docker in your local system

3.Install the docker

4. Login the docker in the power shell by the command

docker login

Pwd

5.After login you need to build the image:

docker build -t predquilty .

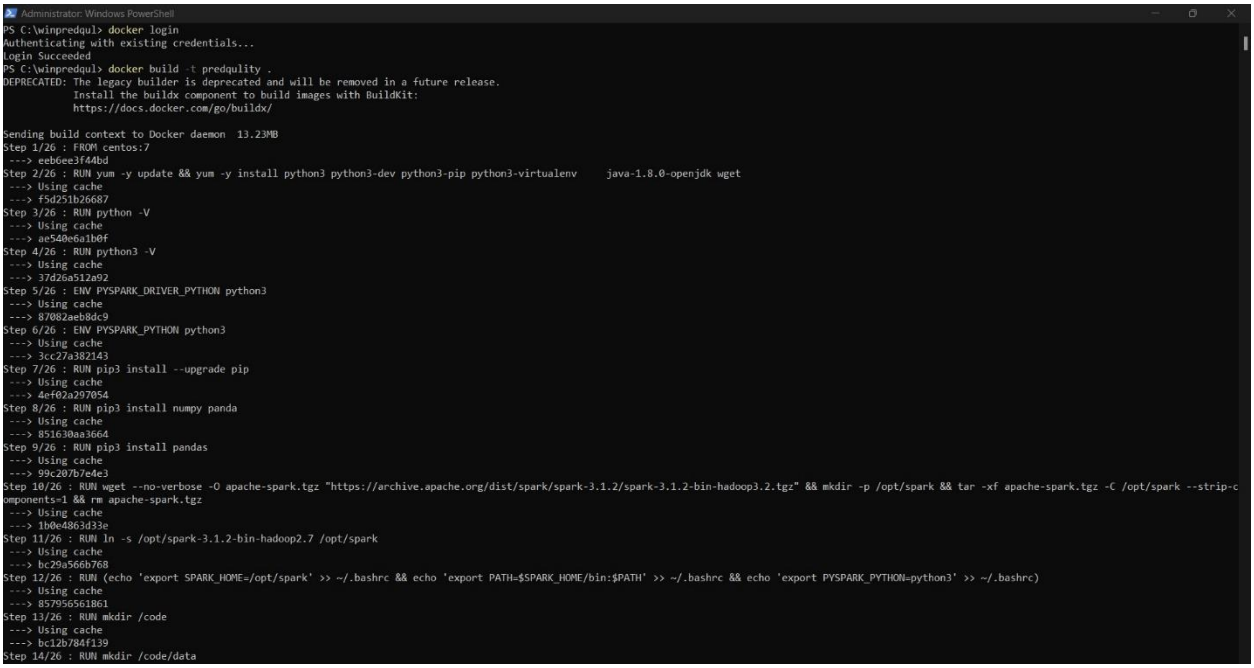

6. The push and pull into the docker hub repository:

PUSH:

docker tag predquility sg2656/ predquility

docker push sg2656/predquility

PULL:

docker pull sg2656/predquility

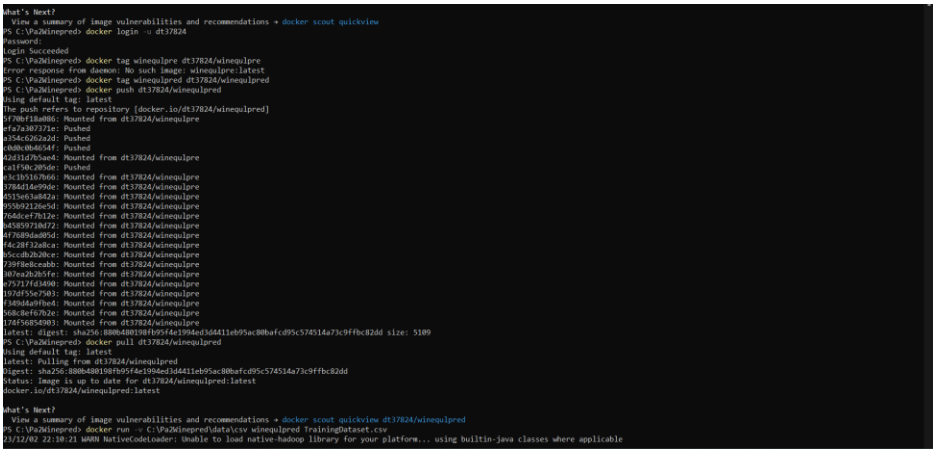

7. Now Your test data file should be kept in a special folder called "dir." Use the Docker container to mount this directory, then run the container by issuing the following command..

docker run -v C:\Pa2 predquility\data\csv Predquility testdata.csv

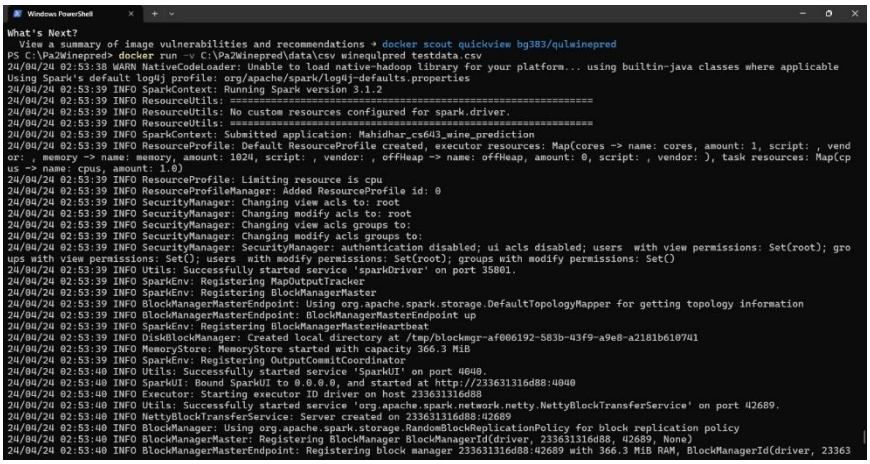

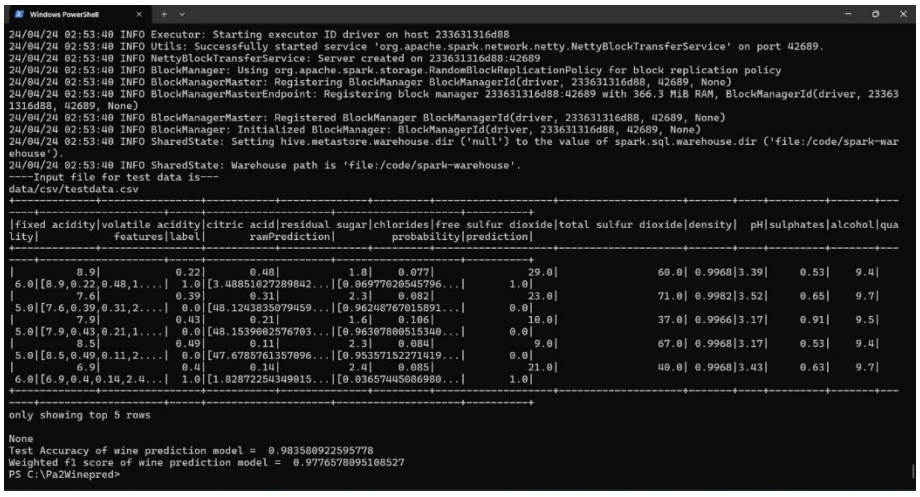

Conclusion: As can be seen in the accompanying graphic, while estimating the wine quality, the accuracy was around 98%.